

Karan Shukla

website: karanshukla.com | email: shuklak13@gmail.com | ph: 214-469-8677

Skills

- ❖ Programming: Python, Java, Scala, R, Git, Bash, SQL
- ❖ Technologies: Deep Learning, Machine Learning, Docker, Kubernetes, Kafka, Hadoop, Spark

Education

- ❖ M.S., Computer Science University of Texas, Dallas Spring 2018
- ❖ B.S., Computer Science [Honors] University of Texas, Dallas Spring 2017
Minor in Statistics Magna Cum Laude (GPA: 3.85) National Merit Phi Kappa Phi

Work Experience

- ❖ Engineering Residency at Google, Mountain View Fall 2018-2019
- ❖ Strategy Consultant Intern at IBM Chief Analytics Office Summer 2017
 - Extracted, cleaned, and transformed real-time text data with Elasticsearch, MongoDB, and Python for use in a text classification and event detection model
 - Led initiative to modernize a web application by...
 - containerizing it with Docker
 - deploying it to Kubernetes and Bluemix/Cloud Foundry clouds
 - creating a REST API for developer access, documented with Swagger
 - speeding up computations with Apache Spark
- ❖ Intern at Citigroup Spring 2017
 - Project manager for a team of five interns designing a virtual reality application in C#
- ❖ Data Science Intern at IBM Extreme Blue Summer 2016
 - Built real-time anomaly detection system using Spark, Kafka, and Scala, using both supervised and unsupervised ML techniques
 - Used Scrum agile framework to bring our idea from conception to prototype to C-suite executive presentation in less than three months with a team of four interns
 - Recovered from disastrous mistakes using Git version control
- ❖ Data Science and Analytics Intern at Verizon Summer 2015
 - Developed statistical programming approaches in R to solve a variety of consumer business problems, including...
 - allocating resources using regression techniques
 - identifying upselling opportunities via classification and segmentation
 - visualizing customer sentiment with ggplot2
 - Queried data with SQL and created visualizations in Excel
 - Learned and utilized a variety of regression techniques, including stepwise regression, LASSO feature selection, and autoregressive time series modelling

Courses

- ❖ Graduate Courses: Database Design Machine Learning Statistics for AI and ML
Computer Networks Fault Tolerance Cloud Computing Distributed Computing
- ❖ Coursera Certifications: [Deep Learning Specialization](#) taught by Professor Andrew Ng

Independent Projects

- ❖ [Tunify](#) won Best Microsoft Hack Aggregate + visualize geo-based Spotify playlists HackTX 2017
- ❖ [anonymizeR](#) Open-source R package for anonymizing data Fall 15 – Spring 16
- ❖ [SentweetSearch](#) Twitter sentiment analysis app built in Python Fall 13 – Spring 14
- ❖ [Final Fantasy](#) Final Fantasy-themed game built with Java Swing Fall 12 – Spring 14

Extracurriculars

- ❖ ACM UTD, Director of Public Relations Spring 17 – Spring 18
- ❖ UTD Student Government, Technology Committee Fall 14 – Spring 15